

OPERATING SYSTEMS PROJECT

REVIEW 2 — METHODOLOGY AND PARTIAL IMPLEMENTATION (LAB ASSESSMENT 9)

Event-Driven Runtime-Adaptive CPU Scheduling

with Sequential Tabular Q-Learning

Single-CPU discrete-event simulator • causal arrived-work state abstraction • sequential tabular Q-learning • seven seeded workload families • held-out crossed-bootstrap evaluation

Field	Value
Project title	Event-Driven Runtime-Adaptive CPU Scheduling with Sequential Q-Learning
Review stage	Review 2 — Methodology and Partial Implementation (Lab Assessment 9)
Repository	github.com/lostpilot404/OS-Project
Language / runtime	Python 3.11.2 (NumPy 2.4.6, pandas 3.0.6, Matplotlib 3.11.2)
Automated verification	273 checks, all passing (re-run during this build)
Dataset	7 synthetic families; 6,000 train / 140 validation / 210 held-out workloads
Report generated	10 October 2026

6 / 6	3,150	6,000	210	162 × 4	273
major modules implemented	processes in the held-out dataset	training workloads	held-out test workloads	causal states × actions	automated tests passing

Summary of the work presented. The proposed system replaces a fixed CPU scheduling policy with a causal, event-driven runtime controller. A single-CPU discrete-event simulator keeps one persistent FIFO ready queue and, at every dispatch or Round-Robin quantum boundary, asks a controller to choose FCFS, SJF, Round Robin or Priority for the next execution segment of the same evolving trace. The controller may observe only work that has already arrived; five causal features are discretised into 162 states and a tabular Q-learning agent with an undiscounted interval reward ($\Sigma r = -\text{total waiting time} / q$) learns when to switch. This report documents the methodology (§1), the architecture (§2), the six modules and their workflows (§3–4), the generated dataset (§5), the tools and algorithms (§6), the experimental plan and metrics (§7), the implementation completed so far with its outputs (§8), and the challenges and remaining plan (§9). Current measured status on the 210 held-out workloads: runtime Q-learning 104.789 mean waiting time versus 215.729 for Round Robin and 93.345 for SJF, with 273/273 automated checks passing.

Contents

The report follows the Review 2 rubric item by item; the nine headings below are the nine required items.

Contents	1
1. Detailed methodology of the proposed system	3
1.1 Research questions	3
1.2 Methodology pipeline	3
1.3 Causal observation contract	3
1.4 State abstraction and learning objective	4
1.5 Methodological controls	4
2. Overall conceptual, architectural and pipeline diagram	5
2.1 Runtime control flow	5
3. Description of the major modules	7
M1 — Workload generator and dataset builder	7
M2 — Discrete-event runtime environment	7
M3 — Causal observation and state encoder	7
M4 — Sequential Q-learning controller	7
M5 — Predeclared heuristic controller	7
M6 — Experiment orchestration, metrics and reporting	8
4. Workflow diagram for each major module	9
M1. Workload generator and dataset builder	9
M2. Discrete-event runtime environment	9
M3. Causal observation and state encoder	10
M4. Sequential Q-learning controller	11
M5. Experiment orchestration and evaluation	12
M6. Metrics and reporting	12
5. Dataset details and data collection procedure	14
5.1 Data collection procedure	14
5.2 Dataset inventory	14
5.3 Family specification	14
5.4 Measured dataset statistics	14
5.5 Data integrity and assumptions	15
6. Tools, technologies, algorithms and frameworks used	16
6.1 Tools and technologies	16
6.2 Algorithms	16
6.3 Controller pseudocode	16
6.4 Frameworks and why none of the heavy ones	16
7. Experimental plan and evaluation metrics	18

7.1 Experimental plan	18
7.2 Evaluation metrics	18
7.3 Statistical procedure	18
7.4 Results obtained with this plan	18
7.5 What the plan has established so far	20
8. Implementation progress	21
8.1 Demonstration 1 — automated verification	21
8.2 Demonstration 2 — end-to-end run output	21
8.3 Demonstration 3 — learned behaviour and artifacts	22
9. Challenges encountered and plan for completing the remaining work	24
9.1 Challenges encountered	24
9.2 Plan for completing the remaining work	24
9.3 Sequence for the remaining work	24

1. Detailed methodology of the proposed system

1.1 Research questions

ID	Question
RQ1	Can a single-CPU scheduler switch between FCFS, SJF, Round Robin and Priority at run time, using only information available at that instant?
RQ2	Does such switching reduce mean waiting time relative to the fixed policies it switches between, under a declared context-switch cost?
RQ3	Does a learned switching policy add value over a hand-written adaptive rule that sees the same observation?
RQ4	Do the gains, if any, survive on a workload family that was never seen in training, and are they larger than the paired bootstrap noise?

1.2 Methodology pipeline

The method is a closed loop between a discrete-event environment that owns the future and a controller that may see only the past and present. Stages 3–7 run once per decision epoch inside a single trace; stage 8 runs once per experiment.

Methodology pipeline (Review 2)




Figure 1.1: Methodology pipeline. Stages 3–7 form the online decision loop; stage 8 is the offline evaluation that produces every number reported here.

Stage	Implemented in	What happens	Output
1. Dataset	workload/generator.py	Produce seven seeded families of 15-process workloads with SHA-256 derived seeds.	RuntimeWorkload list + manifest
2. Simulation	scheduler/runtime.py	Advance one CPU: arrivals, one FIFO ready queue, remaining bursts, switch cost.	decision epochs + schedule trace
3. Observation	scheduler/runtime.py	Freeze an immutable snapshot of arrived/ready work only (causal contract).	RuntimeObservation
4. Encoding	rl/runtime_state.py	Discretise five causal features into 162 tabular states ($3 \times 3 \times 3 \times 3 \times 2$).	state index $s \in [0, 162)$
5. Decision	rl/runtime_controller.py	ϵ -greedy (training) or greedy (evaluation) choice over four policy actions; explicit fallback in unseen states.	action $a \in \{\text{FCFS, SJF, RR, Priority}\}$
6. Execution	scheduler/runtime.py	Run the selected segment (RR bounded by the quantum), charge 1 unit when the running PID changes, settle endpoint arrivals.	advanced simulated time
7. Reward	scheduler/runtime.py	Measure the waiting-time area ΔW accumulated by ready processes; return $r = -\Delta W / q$ ($\gamma = 1$).	interval reward r_t
8. Evaluation	experiments/runtime_experiment.py	Greedy read-only runs on held-out workloads; paired differences with a family-stratified crossed bootstrap.	metric tables + CIs

Table 1.1: The eight methodology stages and the module that owns each.

1.3 Causal observation contract

The environment privately owns the complete event list so it can admit future arrivals at their arrival times; the controller receives an immutable `RuntimeObservation`. This is the methodological core of the project: without it, any measured gain could come from information a real kernel could not have.

Quantity	Visible to the controller?	Methodological reason
Current simulated time	Included	Defines the decision epoch.

Quantity	Visible to the controller?	Methodological reason
Previously selected policy	Included	The action from the immediately preceding epoch (None at the first).
Ready processes (arrived, unfinished)	Included	PID, arrival time, known burst, remaining burst, priority.
Arrived / completed counts	Included	Progress without knowing the workload size.
Mean burst and priority spread over arrived work	Included	Aggregates restricted to processes that have already arrived.
Mean arrival age of ready jobs	Included	current_time – arrival_time; includes prior CPU service, not queue waiting.
Unarrived process data, future bursts/arrivals	Excluded	Would break deployability and fairness of the comparison.
Total workload size, workload identity/family	Excluded	Prevents the controller from memorising generator families.
Counterfactual or completed-schedule metrics	Excluded	No peeking at outcomes of policies that were not run.

Table 1.2: The causal observation contract (scheduler/runtime.py, rl/runtime_state.py). A paired hidden-future test asserts that changing unarrived processes leaves the observation and action prefixes bit-for-bit identical.

1.4 State abstraction and learning objective

#	Feature	Bins	Discretisation (q = 4)	Why it is causal
1	ready_count	3	{1}, {2–3}, {4+}	Instantaneous ready-queue contention.
2	completed_count	3	{0}, {1–3}, {4+}	Trace progress without the workload size.
3	median_ready_remaining_burst	3	$\leq q, \leq 4q, > 4q$	Remaining service demand of ready jobs.
4	mean_ready_arrival_age	3	$\leq q, \leq 4q, > 4q$	Age proxy for ready jobs (includes prior service).
5	arrived_work_priority_spread	2	zero / non-zero	Whether arrived processes have distinct priorities.

Table 1.3: $3 \times 3 \times 3 \times 3 \times 2 = 162$ states; $162 \times 4 = 648$ state-action pairs.

Each trace is one sequential episode. Between consecutive epochs the environment measures the waiting-time area ΔW_t accumulated by ready processes (including any switch interval) and returns a scaled interval reward:

$$r_t = -\Delta W_t / \max(1, q), \gamma = 1.0 \Rightarrow \sum_t r_t = -\text{TotalWaitingTime} / q$$

Because $\gamma = 1$ and the episode is finite, maximising the undiscounted return is identical to minimising total waiting time — the objective is not a proxy for it. The update is

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r_t + \gamma \max_{a' \in \text{visited}(s')} Q(s', a') - Q(s, a)], \alpha = 0.1$$

Terminal transitions omit the bootstrap; the first visit to a state forces a uniform random action; greedy selection and bootstrapping mask unvisited actions; and a wholly unseen evaluation state triggers an explicit, counted fallback to the causal heuristic instead of a silent zero-Q tie.

1.5 Methodological controls

Control	How it is enforced
Fixed-policy baselines	FCFS, SJF, Round Robin and Priority run on the identical workloads and pay the identical switch cost.
Predeclared adaptive rule	A non-RL heuristic (RR $\rightarrow$ Priority $\rightarrow$ SJF $\rightarrow$ FCFS) was fixed before any test result was seen, so learning is not credited for what fixed rules already do.
Blinded held-out split	Final test (master seed 19301) is generated and evaluated after training; the pipeline asserts zero fingerprint overlap with training.
Out-of-family generalisation	poisson_arrivals is excluded from training entirely and used only at evaluation time.
Independent seeds	Five agents (7101–7105) are trained separately; Q-learning results average all five.
Paired uncertainty	Differences are paired per workload and per model seed, then resampled with a family-stratified crossed bootstrap (1,000 replicates).
Determinism	Evaluation is greedy and read-only; artifacts regenerate bit-for-bit from the declared seeds.
Separation of timing	Observation, selection, Q-update and simulator wall time are measured in separate columns and never charged to simulated time.

Table 1.4: Controls that keep the comparison honest.

2. Overall conceptual, architectural and pipeline diagram

The system is layered so that the information boundary is structural, not conventional: the lower (blue) layer owns the future and the accounting, the upper (amber) layer only ever receives a frozen snapshot of arrived work. The controller's reply is a single policy name; it cannot touch the queue, the bursts or the event list.




Figure 2.1: Conceptual architecture. Blue = environment (owns future arrivals, ready queue and switch accounting); amber = controller path (arrived work only); green = the action returned to the environment.

Layer	Components	Responsibility
Data layer	workload/models.py, workload/generator.py	Defines Process/Workload objects and generates the seven seeded families with SHA-256-derived seeds.
Simulation layer	scheduler/runtime.py	Discrete-event loop, FIFO ready queue, remaining bursts, switch cost, schedule trace.
Baseline layer	scheduler/{fcfs,sjf,round_robin,priority}.py	Standalone fixed policies, preserved so runtime actions can be checked against them.
Perception layer	rl/runtime_state.py	Encodes an immutable observation into one of 162 causal states.
Decision layer	rl/runtime_controller.py, rl/runtime_heuristic.py	Learned policy (tabular Q-learning) and the predeclared non-RL rule.
Experiment layer	experiments/runtime_experiment.py, runtime_config.py	Splits, training, validation, held-out testing, paired bootstrap, artifact export.
Evaluation layer	evaluation/metrics.py, evaluation/comparison.py	Metric definitions and paired/bootstrap comparison helpers.

Table 2.1: Architectural layers and what each one owns.

2.1 Runtime control flow

- The environment advances to the next decision epoch (a dispatch point, a completion, or a Round-Robin quantum boundary) and admits every arrival at or before that instant in (arrival_time, pid) order.
- It freezes an observation and passes it to the controller; the controller returns one of four policy names.
- If the selected process differs from the running PID, one unit of switch cost is charged before the segment starts; arrivals during that interval are queued but the dispatch already chosen is not reconsidered.
- The segment runs (Round Robin bounded by the quantum, others to completion); endpoint arrivals are admitted first, then an unfinished Round-Robin process is appended to the tail.
- The environment measures the waiting-time area accumulated by ready processes, converts it to the interval reward, and hands it back to the controller for the Q update.

- A switch changes only the next selection rule and service length — the queue, the remaining bursts and the switch accounting are preserved across switches.

3. Description of the major modules

Six modules carry the whole system: two build and run the simulated machine, two implement the competing controllers, one is the non-learning reference rule, and one orchestrates the experiment and its statistics. Each is described below with its contract; §4 gives the workflow diagram for each.

ID	Module	Files	Input → output	Responsibility	Status
M1	Workload generator and dataset builder	workload/generator.py, workload/models.py	family config + seed → validated Workload with fingerprint	Generates the seven synthetic families, validates every process (positive burst, non-negative arrival, priority bounds) and records split/seed/fingerprint metadata.	Implemented
M2	Discrete-event runtime environment	scheduler/runtime.py	Workload + controller + (quantum, switch cost) → schedule trace and six metrics	Owens the private event list, the single FIFO ready queue, remaining bursts, switch-cost accounting, the causal observation and the interval reward.	Implemented
M3	Causal observation and state encoder	rl/runtime_state.py	RuntimeObservation → state index in [0, 162)	Extracts five causal features, discretises them relative to the quantum and flattens them in mixed-radix order. Accepts an observation only — never a Workload.	Implemented
M4	Sequential Q-learning controller	rl/runtime_controller.py	observation (encode) / reward (update) → policy action; Q table and visit counts	Tabular Q-learning over 648 pairs: ϵ -greedy exploration, first-visit forcing, unvisited-action masking, read-only greedy evaluation and a counted heuristic fallback.	Implemented
M5	Predeclared heuristic controller	rl/runtime_heuristic.py	RuntimeObservation → policy action	Non-RL adaptive rule (RR → Priority → SJF → FCFS) sharing the same observation contract, declared before any test result was inspected.	Implemented
M6	Experiment orchestration, metrics and reporting	experiments/runtime_experiment.py, evaluation/metrics.py	configuration + trained models → CSV/JSON artifacts, paired CIs, markdown report	Builds disjoint splits, trains five agents, validates, runs the held-out test, computes paired bootstrap intervals and writes every artifact.	Implemented

Table 3.1: Major modules, contracts and current status.

M1 — Workload generator and dataset builder

Seven frozen family configurations (burst distribution and range, arrival pattern, priority pattern) are turned into workloads of 15 processes. Each workload seed is derived with SHA-256 from (master seed, split tag, family index, repetition), so the whole dataset is reproducible from three integers and no external data is required. The generator validates every process and the pipeline records a workload fingerprint for split-disjointness checks.

M2 — Discrete-event runtime environment

The simulator is deliberately the only component that knows the future. It advances time to the next epoch, maintains one persistent FIFO ready queue, tracks remaining bursts, charges one unit whenever the running PID changes, and produces the schedule trace from which all metrics are computed. Running it with a constant action reproduces the standalone FCFS, SJF, Round Robin and Priority implementations exactly — a property that is asserted by tests, not merely claimed.

M3 — Causal observation and state encoder

The encoder's API accepts a RuntimeObservation and nothing else, which makes the information boundary checkable by inspection. Five features are discretised into $3 \times 3 \times 3 \times 3 \times 2 = 162$ states. Bin edges are expressed relative to the quantum so the abstraction stays meaningful if the quantum changes.

M4 — Sequential Q-learning controller

A tabular agent over 162×4 pairs. During training it follows an episode-level ϵ schedule, forces a random action on a state's first visit and masks unvisited actions in the greedy step and in the bootstrap. During evaluation it is greedy and read-only; an unseen state produces an explicit fallback to M5, and the number of fallbacks is reported rather than hidden.

M5 — Predeclared heuristic controller

Round Robin if at least three jobs are ready and their mean arrival age is at least one quantum; else Priority if at least two ready jobs differ in priority; else SJF if the largest visible remaining burst is at least twice the smallest; else FCFS. It

exists so that the learned controller is compared against adaptivity itself, not only against fixed policies.

M6 — Experiment orchestration, metrics and reporting

Builds the three splits, asserts they are fingerprint-disjoint, trains the five agents, evaluates greedily on validation and on the held-out test, computes paired differences with a family-stratified crossed bootstrap, and writes every metric table, Q value, visit count, coverage figure and the illustrative trace to disk.

4. Workflow diagram for each major module

Each module below is shown as its own workflow: the numbered steps are the order of execution inside that module, and the arrows are the data passed between them.

M1. Workload generator and dataset builder

workload/generator.py → workload/models.py




Figure 4.1: Workflow of Workload generator and dataset builder.

#	Step	Detail
1	Select family configuration	7 frozen parameter sets (burst, arrival, priority)
2	Derive the workload seed	SHA-256 (master seed, split tag, family index, repetition)
3	Sample 15 processes	burst, arrival time and priority drawn per family rules
4	Validate the workload	positive bursts, ordered arrivals, priority bounds
5	Fingerprint and tag	SHA-256 workload fingerprint + split/repetition metadata
6	Assemble the split	6,000 train / 140 validation / 210 held-out workloads

M2. Discrete-event runtime environment

scheduler/runtime.py – one iteration per decision epoch




Figure 4.2: Workflow of Discrete-event runtime environment.

#	Step	Detail
1	Advance to the next epoch	dispatch point, completion or Round-Robin quantum boundary
2	Admit arrivals	all arrivals at or before the epoch, in (arrival_time, pid) order
3	Freeze the observation	arrived/ready work only — the future stays private
4	Ask the controller for an action	FCFS, SJF, Round Robin or Priority
5	Charge the switch cost	1 unit if the running PID changes; first dispatch is free
6	Execute the segment	min(quantum, remaining) for RR, otherwise to completion
7	Settle the endpoint	admit arrivals, then RR tail re-enqueue or mark completed
8	Return the interval reward	$r = -\Delta W / q$, measured by the environment

M3. Causal observation and state encoder

rl/runtime_state.py




Figure 4.3: Workflow of Causal observation and state encoder.

#	Step	Detail
1	Receive the observation	immutable RuntimeObservation (no Workload object)
2	Extract five causal features	ready count, completed count, median remaining burst, mean arrival age, priority spread
3	Discretise each feature	bin edges expressed relative to the quantum $q = 4$
4	Flatten in mixed-radix order	$3 \times 3 \times 3 \times 3 \times 2 \rightarrow$ one integer
5	Return the state index	$s \in [0, 162)$; consumed by the controller only

M4. Sequential Q-learning controller

rl/runtime_controller.py




Figure 4.4: Workflow of Sequential Q-learning controller.

#	Step	Detail
1	Encode the observation	$s = \text{encoder}(\text{observation})$
2	Select an action	training: ϵ -greedy + forced random on first visit; evaluation: greedy over visited actions
3	Handle unseen states	explicit fallback to the causal heuristic (counted and reported)
4	Receive the reward	$r = -\Delta W / q$ from the environment
5	Update the table	$Q(s,a) += \alpha(r + \gamma \max_{\text{visited}} Q(s',.) - Q(s,a))$; terminal omits the bootstrap
6	Persist Q values and visits	written to runtime_q_table.csv per training seed

M5. Experiment orchestration and evaluation

experiments/runtime_experiment.py



Figure 4.5: Workflow of Experiment orchestration and evaluation.

#	Step	Detail
1	Build the three splits	independent seeded streams per split
2	Assert split disjointness	zero SHA-256 fingerprint overlap (fails loudly otherwise)
3	Train five agents	1,200 sequential episodes per seed, six families
4	Validate	140 workloads — reporting and demo selection only, never tuning
5	Run the held-out test	210 untouched workloads, greedy and read-only
6	Compare pairwise	paired differences + 1,000-replicate crossed bootstrap
7	Write the artifacts	CSV/JSON tables, Q table, coverage, demo trace, report

M6. Metrics and reporting

evaluation/metrics.py, evaluation/comparison.py




Figure 4.6: Workflow of Metrics and reporting.

#	Step	Detail
1	Collect the schedule trace	start/end time per (pid, segment)
2	Derive per-process times	waiting, turnaround and response per process
3	Aggregate per workload	means, CPU utilisation, throughput, context switches
4	Aggregate per method and family	mean $\pm$ SD across workloads (and model seeds)
5	Compute paired differences	target – reference per workload, per seed for Q-learning
6	Render the outputs	CSV tables, summary JSON, markdown report, figures

5. Dataset details and data collection procedure

5.1 Data collection procedure

No external or personal data is used. The dataset is synthetic and generated entirely by the project's own generator, so it is reproducible from three integers and can be regenerated on any machine. The procedure is:

Step	Action	Detail
D1	Fix seven family configurations	burst distribution and range, arrival pattern/window/rate, priority pattern — declared once in config.py and frozen in the experiment configuration.
D2	Derive one seed per workload	workload_seed = SHA-256(master seed, split tag, family index, repetition); no global RNG state is shared between workloads.
D3	Generate 15 processes per workload	burst, arrival time and priority are drawn per the family rules, then validated (positive burst, non-negative arrival, priority within bounds).
D4	Assign workloads to splits	6,000 training workloads (6 families × 1,200 episodes × 5 seeds), 140 validation (7 × 20) and 210 held-out test (7 × 30).
D5	Fingerprint and check disjointness	each workload gets a SHA-256 fingerprint; the pipeline asserts zero overlap between training, validation and test.
D6	Store only what is needed	metrics, manifest (split/seed/fingerprint/process count) and compact summaries are committed; the large per-decision event logs are regenerated by the CLI and gitignored.
D7	Regenerate on demand	`python main.py runtime-experiment` rebuilds every workload and every metric from the declared seeds.

Table 5.1: Data collection procedure.

5.2 Dataset inventory

Split	Families	Repetitions	Workloads	Processes	Master seed
Training (sequential episodes)	6 families	1,200 episodes × 5 seeds	6,000	90,000	7101–7105
Validation	7 families	20 repetitions	140	2,100	8201
Final held-out test	7 families	30 repetitions	210	3,150	19301

Table 5.2: Dataset inventory. Every workload has exactly 15 processes; poisson_arrivals is excluded from training and used only at evaluation time.

5.3 Family specification

Family	Description	Burst range	Burst law	Arrivals	Priority
short_jobs	Short-job-heavy: 80% of bursts drawn from 1-5 time units.	1–50	bimodal	uniform	1–5
long_jobs	Long-job-heavy: 20% of bursts drawn from 1-5 time units.	1–50	bimodal	uniform	1–5
mixed	Uniform mix of bursts over the full 1-50 range.	1–50	uniform	uniform	1–5
cpu_bursty	Long CPU bursts (20-50) with dense uniform arrivals.	20–50	uniform	uniform	1–5
priority_skewed	Uniform bursts; 70% of priorities lie in the 1-2 band.	1–50	uniform	uniform	1–5
staggered_interactive	One long process at time zero plus 14 short processes arriving at integer gaps of 4-6 time units.	25–50	staggered_interactive	staggered	1–1
poisson_arrivals	Held-out evaluation condition: Poisson arrivals (rate 0.5) and uniform bursts from 1-20.	1–20	uniform	poisson	1–5

Table 5.3: The seven declared workload families (15 processes each).

5.4 Measured dataset statistics

Family	Workloads	Processes	Burst (mean ± SD)	Mean arrival	Mean total work	High-priority share
short jobs	30	450	8.5 ± 12.1	9.6	128	40%
long jobs	30	450	22.2 ± 15.4	9.6	333	41%
mixed	30	450	24.7 ± 14.3	10.0	370	41%
cpu bursty	30	450	35.0 ± 9.0	7.8	525	38%
priority skewed	30	450	26.3 ± 14.6	10.0	395	82%
staggered interactive	30	450	4.3 ± 9.1	35.4	65	100%
poisson arrivals	30	450	10.7 ± 5.6	13.5	160	44%
All families	210	3150	18.8 ± 15.8	13.7	282	55%

Table 5.4: Statistics of the 210 held-out workloads (3,150 processes), recomputed with the project's generator at report-build time — these are properties of the actual dataset, not of the configuration alone. High priority means priority ≤ 2 .

Generated dataset profile — 210 held-out workloads x 15 processes = 3,150 processes (rebuilt from the declared seeds)




Figure 5.1: Profile of the generated held-out dataset: burst-time and arrival-time spread per family, total work per workload and high-priority share.

5.5 Data integrity and assumptions

- **Reproducibility:** seeds are derived with SHA-256, so a workload is a pure function of (master seed, split tag, family index, repetition).
- **Disjointness:** the pipeline fails loudly if any fingerprint appears in two splits (current run: overlap = false, 6,000 distinct training fingerprints).
- **Burst-knowledge assumption:** exact burst lengths become known when a process arrives. This is the assumption already required by standalone SJF; no burst predictor is included.
- **Storage discipline:** per-decision event logs are large and are regenerated by the CLI, so they are gitignored; the manifest, metrics, Q tables and summaries are committed.
- **No personal or licensed data:** the dataset is entirely synthetic, so there are no privacy or licensing constraints.

6. Tools, technologies, algorithms and frameworks used

6.1 Tools and technologies

Category	Item	Role in the project
Language	Python 3.11.2	Whole implementation, CLI and report generator
Numerics	NumPy 2.4.6	Q table as an array, bootstrap resampling
Data handling	pandas 3.0.6	Metric aggregation and CSV artifact export
Plotting	Matplotlib 3.11.2	All figures in this report
Testing	pytest 9.1.1	273 automated checks
PDF generation	ReportLab 5.0.1	Renders this report from the committed artifacts
Hashing	hashlib (SHA-256, stdlib)	Seed derivation and workload fingerprints
Standard library	argparse, dataclasses, json, pathlib	Configuration objects, CLI, artifact serialisation
Version control	Git / GitHub	Branch-per-change workflow, pull-request review
Environment	Linux 6.1 (x86-64), venv + pip	Development, experiment runs and report builds

Table 6.1: Tools and technologies. No deep-learning framework is used — the controller is deliberately tabular so it stays auditable.

6.2 Algorithms

Algorithm	Definition	Where it is used
FCFS (action 0)	Dispatch the head of the ready queue, non-preemptively.	Baseline and runtime action; after an RR requeue it follows live FIFO order.
SJF (action 1)	Smallest remaining burst first, non-preemptive, ties by ready order.	Baseline and runtime action; optimal for mean waiting time under known bursts.
Round Robin (action 2)	Head of the queue for min(quantum, remaining); unfinished jobs go to the tail after endpoint arrivals.	Baseline and runtime action; bounds response time.
Priority (action 3)	Highest static priority first, non-preemptive, ties by ready order.	Baseline and runtime action; lower numeric value = higher priority.
Tabular Q-learning	$Q(s,a) \leftarrow \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$ with ϵ -greedy exploration and unvisited-action masking.	The learned runtime controller (162 states $\times$ 4 actions).
Predeclared heuristic	RR $\rightarrow$ Priority $\rightarrow$ SJF $\rightarrow$ FCFS cascade over the same observation.	Non-learning adaptive reference.
Crossed bootstrap	Resample workloads within family and model seeds independently; 1,000 replicates.	Paired 95% confidence intervals for every comparison.
Discrete-event simulation	Event-list advancement with a persistent FIFO ready queue and switch-cost charging.	The machine being controlled.

Table 6.2: Algorithms implemented and their role.

6.3 Controller pseudocode

```

SELECT_ACTION(obs):
    s <- encode(obs)                # training: epsilon-greedy; evaluation: greedy
    if evaluating:
        if no visited action in s:   # wholly unseen state
            fallback_count += 1
            return HEURISTIC(obs)    # predeclared rule, reported not hidden
        return argmax over visited actions of Q[s, a]
    if first_visit(s) or rand() < epsilon(episode):
        return uniform_random_action()
    return argmax over visited actions of Q[s, a]

UPDATE(s, a, r, s', terminal):
    target <- r if terminal else r + gamma * max over visited actions of Q[s', .]
    Q[s, a] <- Q[s, a] + alpha * (target - Q[s, a])
    visits[s, a] <- visits[s, a] + 1

HEURISTIC(obs):
    # predeclared, no learning
    if ready_count >= 3 and mean_arrival_age >= quantum: return Round Robin
    if >= 2 ready jobs have distinct priorities:         return Priority
    if max(remaining_burst) >= 2 * min(remaining_burst): return SJF
    return FCFS

```

6.4 Frameworks and why none of the heavy ones

The controller is a 162×4 table rather than a neural network, so no deep-learning framework is needed. This is a deliberate methodological choice: every Q value, visit count and fallback can be inspected directly in the committed artifacts, which is what makes claims about state coverage and unseen-state behaviour verifiable instead of anecdotal.

7. Experimental plan and evaluation metrics

7.1 Experimental plan

Step	Action	Detail
P1	Fix the design before looking at results	Families, splits, seeds, quantum (4), switch cost (1), learning rates and the heuristic rule are all declared in configuration and committed.
P2	Train five independent agents	Seeds 7101–7105; 1,200 sequential episodes each on six families (poisson_arrivals withheld).
P3	Validate once	140 workloads (seed 8201) used only for reporting and for the predeclared illustrative-demo rule — never for tuning.
P4	Test once, held out	210 untouched workloads (seed 19301); the fixed policies and the heuristic are deterministic (one run each), Q-learning is evaluated by all five models greedily and read-only.
P5	Compare every method against every other	Paired per-workload differences, plus independent resampling of the five model seeds.
P6	Quantify uncertainty	Family-stratified crossed bootstrap, 1,000 replicates, seed 104729; 95% intervals on every headline comparison.
P7	Publish artifacts	Metric tables, Q values, visit counts, coverage, manifest, demo trace and summary JSON are written to results/runtime/ and regenerated by one command.

Table 7.1: The experimental plan, in execution order.

7.2 Evaluation metrics

Metric	Definition	Why it is measured
Mean waiting time	mean over processes of (first service start – arrival)	Primary metric — it is exactly the quantity the reward minimises.
Mean turnaround time	mean over processes of (completion – arrival)	End-to-end latency experienced by a job.
Mean response time	mean over processes of (first dispatch – arrival)	Interactivity; this is what Round Robin buys.
CPU utilisation (%)	busy time / total elapsed time × 100	Cost of switching and idling.
Throughput	processes completed per unit simulated time	Work completed per unit time.
Context switches	number of running-PID changes	Overhead proxy; switch cost is charged for every method.
Policy switches per trace	decisions where the chosen policy differs from the previous	How much the controller actually adapts.
State-action coverage	visited (s, a) pairs / 648, unseen-state fallbacks	Whether the learned table is actually populated where it is used.

Table 7.2: Evaluation metrics.

7.3 Statistical procedure

Comparisons are paired: for each workload the difference (target – reference) is computed on the same workload, and for Q-learning the model seed is resampled independently from the workloads (a crossed bootstrap). Resampling is stratified by workload family, and 1,000 replicates are drawn with seed 104729. An interval that excludes zero on the favourable side is treated as a real difference on this benchmark; an interval that contains zero is reported as inconclusive rather than as a win.

7.4 Results obtained with this plan

Method (210 held-out workloads)	Mean waiting	Mean turnaround	Mean response	CPU util %	Throughput	Ctx switches	Policy switches
FCFS	127.797 ± 79.81	146.612 ± 90.51	127.797 ± 79.81	92.24 ± 6.15	0.077 ± 0.056	14.00 ± 0.00	0.00 ± 0.00
SJF	93.345 ± 68.78	112.161 ± 79.39	93.345 ± 68.78	92.24 ± 6.15	0.077 ± 0.056	14.00 ± 0.00	0.00 ± 0.00
Round Robin	215.729 ± 172.82	234.544 ± 183.45	26.220 ± 9.15	78.22 ± 1.90	0.068 ± 0.054	74.77 ± 40.70	0.00 ± 0.00
Priority	129.207 ± 80.81	148.022 ± 91.53	129.207 ± 80.81	92.24 ± 6.15	0.077 ± 0.056	14.00 ± 0.00	0.00 ± 0.00
Causal heuristic	203.404 ± 159.23	222.220 ± 169.85	40.070 ± 17.05	80.48 ± 2.22	0.071 ± 0.058	67.37 ± 39.49	3.21 ± 0.67
Runtime Q-learning	104.789 ± 71.43	123.604 ± 81.86	86.341 ± 59.82	90.84 ± 7.04	0.076 ± 0.054	17.32 ± 4.66	4.73 ± 2.83

Table 7.3: Held-out final test (master seed 19301, quantum 4, switch cost 1). Mean ± SD across the 210 workloads; Q-learning also averages the five model seeds.

Held-out final test: 210 workloads (seed 19301), 5 Q-learning model seeds




Figure 7.1: Six metrics on the held-out test set (error bars = SD across workloads).

5 independent agents x 1,200 sequential episodes (6,000 training workloads)




Figure 7.2: Sequential training: rolling mean waiting time and episodic return (100-episode window) for each of the five agents.

Target	Reference	Metric	Mean difference	95% bootstrap CI	Reading
Runtime Q	SJF	waiting time	+11.443	[+5.022, +19.297]	favours reference
Runtime Q	SJF	response time	-7.004	[-9.541, -4.305]	favours target
Runtime Q	SJF	context switches	+3.322	[+2.209, +4.555]	favours reference
Runtime Q	SJF	CPU util %	-1.402	[-1.843, -1.005]	favours reference
Runtime Q	Heuristic	waiting time	-98.616	[-106.097, -90.192]	favours target
Runtime Q	Heuristic	response time	+46.272	[+43.472, +49.100]	favours reference
Runtime Q	Heuristic	context switches	-50.050	[-51.929, -48.066]	favours target
Runtime Q	Heuristic	CPU util %	+10.354	[+9.890, +10.834]	favours target
RR	SJF	waiting time	+122.383	[+117.619, +127.315]	favours reference
RR	SJF	response time	-67.126	[-70.219, -64.285]	favours target
RR	SJF	context switches	+60.767	[+58.995, +62.491]	favours reference
RR	SJF	CPU util %	-14.021	[-14.241, -13.760]	favours reference
Heuristic	SJF	waiting time	+110.059	[+106.219, +114.032]	favours reference
Heuristic	SJF	response time	-53.276	[-56.026, -50.873]	favours target

Target	Reference	Metric	Mean difference	95% bootstrap CI	Reading
Heuristic	SJF	context switches	+53.371	[+51.666, +55.191]	favours reference
Heuristic	SJF	CPU util %	-11.756	[-11.992, -11.518]	favours reference
FCFS	SJF	waiting time	+34.451	[+32.640, +36.540]	favours reference
FCFS	SJF	response time	+34.451	[+32.605, +36.440]	favours reference
FCFS	SJF	context switches	+0.000	[+0.000, +0.000]	no difference
FCFS	SJF	CPU util %	+0.000	[+0.000, +0.000]	no difference

Table 7.4: Paired differences (target – reference) with family-stratified crossed-bootstrap intervals.




Figure 7.3: Forest plot of the paired comparisons (green favours the target, red favours the reference).

7.5 What the plan has established so far

- **Against Round Robin** the learned controller is clearly better on waiting time: -110.940 time units, obtained by differencing the two committed comparisons against SJF ($Q - SJF +11.443$ minus $RR - SJF +122.383$; both intervals exclude zero and are far apart, so the ordering is unambiguous). FCFS and Priority sit between the two, at +34.451 and +35.862 relative to SJF.
- **Against the predeclared adaptive rule** it reduces waiting time by -98.616 (95% CI [-106.097, -90.192]) — so the gain comes from learning, not merely from being adaptive.
- **Against SJF it is mixed, and the report says so:** waiting time is worse by +11.443 (95% CI [+5.022, +19.297]) while response time is better by -7.004 (95% CI [-9.541, -4.305]). Non-preemptive SJF with exact bursts and no mid-job switch cost is a very strong reference on this benchmark.
- On the held-out poisson_arrivals family the controller still lands between SJF and Round Robin, so the behaviour is not family memorisation.

8. Implementation progress

All six major modules are implemented, exercised by the automated suite and demonstrated with the outputs below. The requirement for Review 2 is that at least 50% of the proposed work is complete and that the completed modules are demonstrated with valid outputs; the table below records completion per module and the artifact that evidences it.

ID	Module	Completion	Evidence artifact	Status
M1	Workload generator and dataset builder	100%	6,350-row manifest + regenerated dataset statistics (§5.4)	Complete
M2	Discrete-event runtime environment	100%	fixed-action equivalence tests; schedule traces in final-test metrics	Complete
M3	Causal observation and state encoder	100%	162-state table; hidden-future invariance tests	Complete
M4	Sequential Q-learning controller	100%	runtime_q_table.csv (3,240 Q values), coverage report, fallback counts	Complete
M5	Predeclared heuristic controller	100%	heuristic rows in every metric table	Complete
M6	Experiment orchestration, metrics and reporting	100%	final-test, validation, paired-bootstrap and coverage tables; generated report	Complete
—	Review 2 report generator (this document)	100%	tools/build_review2_report.py	Complete

Table 8.1: Implementation status. Overall, the proposed Review 2 scope (methodology, simulator, controller, experiment, evaluation) is complete; the items listed in §9 are extensions beyond the proposed scope.

8.1 Demonstration 1 — automated verification

Test module	Checks	What it protects
test_adaptive.py	18	Legacy offline selector: state encoding, decisions, read-only evaluation
test_experiment_design.py	4	Frozen seeds, split sizes and configuration invariants
test_experiments.py	30	Legacy experiment pipeline, metrics and report generation
test_fcfs.py	11	FCFS ordering, idle handling and metric computation
test_generator.py	7	Workload family generation and reproducibility from seeds
test_metrics.py	14	Waiting / turnaround / response / utilisation / throughput definitions
test_priority.py	14	Priority ordering, ties and non-preemption
test_q_learning.py	29	Q updates, epsilon schedule, masking, visit counts, evaluation
test_reward.py	16	Reward composition and clipping against the reference policies
test_round_robin.py	16	Quantum boundaries, re-enqueue order and completion behaviour
test_runtime_experiment.py	4	Split disjointness, determinism, artifact regeneration
test_runtime_scheduler.py	23	Runtime environment: fixed-action equivalence, causality, reward identity
test_schedule_result.py	14	Schedule trace validation and aggregate result objects
test_sjf.py	12	SJF selection, tie-breaking and non-preemption
test_state.py	24	Offline state discretisation and mixed-radix indexing
test_workload_models.py	37	Process/workload models, validation and trace invariants

Table 8.2: Test inventory, 273 checks collected and executed while building this report.

```
$ python -m pytest -q
..... [ 79%]
..... [100%]
273 passed in 9.58s
```

8.2 Demonstration 2 — end-to-end run output

The full pipeline (training, validation, held-out testing, artifact export) is reproduced by one command. The output below is from a verification run performed on the reporting host and matches the committed artifacts exactly.

```
$ python main.py runtime-experiment --results-dir /tmp/os-project-runtime
Causal event-driven runtime experiment
-----
training seeds           : 7101, 7102, 7103, 7104, 7105
training episodes/model  : 1200
validation workloads     : 140
final held-out test workloads : 210
final-test mean waiting time:
  FCFS                   : 127.797
  SJF                     : 93.345
  Round Robin            : 215.729
  Priority                 : 129.207
  Causal heuristic       : 203.404
  Runtime Q-learning     : 104.789
learned within-trace switches : 15 (23 decisions; runtime_learned_switch_demo.json)
artifacts:
  training_metrics       : runtime_training_metrics.csv
  validation_metrics     : runtime_validation_metrics.csv
  final_test_metrics     : runtime_final_test_metrics.csv
  paired_comparisons     : runtime_paired_comparisons.csv
  state_action_coverage  : runtime_state_action_coverage.csv
  q_table                : runtime_q_table.csv
  workload_manifest      : runtime_workload_manifest.csv
  learned_switch_demo    : runtime_learned_switch_demo.json
  summary                : runtime_summary.json
  report                 : runtime_report.md
```

8.3 Demonstration 3 — learned behaviour and artifacts




Figure 8.1: Learned state-action visitation pooled over the five agents (left) and per-seed coverage (right). Between 267 and 276 of the 648 pairs are visited.

The controller does switch policies inside a single trace. The example below is taken from the validation split by the predeclared rule (training seed 7103, family staggered_interactive, fingerprint 275409d7458331f6): 15 switches over 23 decisions, with mean waiting time 5.933 and mean response time 2.933. It is an illustration of the mechanism, not a performance claim.




Figure 8.2: Illustrative trace. Bars are execution segments coloured by the policy chosen for them; dotted lines mark the instants where the policy changed.

Artifact	Size	Contents
results/runtime/runtime_final_test_metrics.csv	2,100 rows	Six metrics, switch/decision counts and timings on the held-out test
results/runtime/runtime_validation_metrics.csv	1,400 rows	Same columns for the validation split
results/runtime/runtime_training_metrics.csv	6,000 rows	Per-episode sequential training outcomes
results/runtime/runtime_paired_comparisons.csv	36 rows	Paired differences with 95% crossed-bootstrap intervals
results/runtime/runtime_state_action_coverage.csv	15 rows	Visited states/pairs, fallback counts and timings per seed
results/runtime/runtime_q_table.csv	3,240 rows	Q values and visit counts per training seed
results/runtime/runtime_workload_manifest.csv	6,350 rows	Split, seeds and workload fingerprints
results/runtime/runtime_learned_switch_demo.json	1 trace	Observations, actions, switch times and execution trace
results/runtime/runtime_summary.json	—	Configuration, split fingerprints, all means and software versions
figures/review2/*.png	11 figures	Every figure in this report

Table 8.3: Artifacts produced by the completed modules.

9. Challenges encountered and plan for completing the remaining work

9.1 Challenges encountered

Challenge	Why it mattered	How it was resolved
Keeping the controller causal	An earlier whole-workload selector saw future arrivals and counterfactual schedules, so its results could not be interpreted as a runtime controller.	Introduced an immutable RuntimeObservation, made the encoder accept only that object, and added a paired hidden-future test that fails if unarrived changes leak into the observation or action prefix.
Round-Robin and FCFS interaction after a requeue	Once RR appends a preempted job to the tail, 'FCFS' is ambiguous: original arrival order or live queue order?	Defined one persistent FIFO ready queue with canonical (arrival_time, pid) insertion, and specifies that FCFS follows live queue order while SJF/Priority ties preserve it. Fixed-action runs are tested against the standalone schedulers.
Charging switch cost fairly	If the learner pays for switching but baselines do not, the comparison is meaningless; if nothing is charged, switching is free and unrealistic.	One unit is charged whenever the running PID changes, for every method including the fixed policies; arrivals during the switch interval are queued but the already-chosen dispatch is not reconsidered.
Designing a reward that is the objective	A shaped reward can drift away from the metric being reported.	Used $r = -\Delta W/q$ with $\gamma = 1$, so the episodic return is exactly $-\text{total waiting time} / q$; a test asserts the identity.
Sparse coverage of the tabular space	Only ~42% of the 648 state-action pairs are ever visited, and unseen states appear during evaluation.	Kept the abstraction small, forced exploration on first visit, masked unvisited actions in greedy steps and bootstraps, and made unseen states fall back to the predeclared heuristic with the fallback count reported (1–3 decisions per model on the final test).
Honest reporting of a negative result	The learner does not beat non-preemptive SJF on waiting time; it would be easy to bury this.	Reported it as a headline finding with its bootstrap interval, and stated the structural reason (SJF gets exact bursts at arrival and pays no mid-job switch cost).
Host-dependent timing noise	Wall-clock columns make artifacts look non-deterministic across machines.	Observation, selection, Q-update and simulator time are stored in separate columns, never charged to simulated time, and excluded from all schedule comparisons.
Large per-decision event logs	Full causal decision logs are large and would bloat the repository.	They are written by the CLI but gitignored and regenerable from the recorded configuration and seeds; compact metrics, Q tables and manifests are committed.

Table 9.1: Challenges and their resolutions.

9.2 Plan for completing the remaining work

ID	Task	Gap it closes	Planned approach	Deliverable
W1	Sensitivity study	All results are for $q = 4$ and switch cost 1.	Re-run the plan over a grid of quanta and switch costs and report where the ordering changes.	Metric tables per (q, cost) cell
W2	Burst prediction	Exact bursts are assumed known at arrival.	Add an explicit predictor with measured error and re-run the controller under that noise.	Results with prediction error bands
W3	Richer state abstraction	Only ~42% of pairs are visited; 162 states is coarse.	Tile coding or a small function approximator, keeping the same causal observation contract and fallback reporting.	Coverage and metric comparison vs the tabular agent
W4	Multi-objective reward	Waiting time is optimised; response time is not.	Add a constrained/weighted objective and measure the waiting-response trade-off curve.	Pareto front over the two metrics
W5	I/O and blocking	Processes never block, so the queue is CPU-bound only.	Model I/O bursts and blocked/wake transitions in the environment.	Extended environment plus new baselines
W6	Multicore extension	Single CPU only.	Extend to N cores with per-core queues and migration cost; the controller chooses a policy per core.	Multicore metrics and overhead analysis
W7	Trace-driven validation	All workloads are synthetic.	Replay a public scheduling trace through the same environment and report the gap to the synthetic results.	Trace-driven result table

Table 9.2: Remaining work, planned beyond the proposed Review 2 scope.

9.3 Sequence for the remaining work

- **First:** W1 (sensitivity) and W2 (burst prediction) — they test whether the current conclusions are an artefact of fixed parameters or an idealised assumption.
- **Second:** W3 and W4 — they attack the two measured weaknesses directly (coarse coverage, single objective).

- **Third:** W5 and W6 — environment extensions that make the simulator more realistic, each with re-baselined comparisons.
- **Finally:** W7 — external validation on a real trace, which is the strongest available check that the synthetic ordering means anything.